



# Common Developer Crypto Mistakes (with illustrations in Java)

Kevin W. Wall

***BSidesCMH***

March 1, 2019

Copyright © 2017-2019 – Kevin W. Wall – All Rights Reserved.  
Released under Creative Commons Attribution-Noncommercial-Share Alike 3.0  
United States License  
as specified at  
<http://creativecommons.org/licenses/by-nc-sa/3.0/us/>



# Obligatory “It's all about me” page



- 35+ years developer experience, 15+ yrs security experience
  - 17 yrs at (now Nokia) Bell Labs; left as DMTS
  - 3.5 yrs as independent contractor (C++ & Java)
  - 14 years AppSec & InfoSec experience at CenturyLink / Qwest
- Currently: Information Security Engineer at Wells Fargo on Secure Code Review team (5+ yrs)
- OWASP ESAPI for Java
  - Project co-leader
  - Cryptography developer (since Aug 2009)
- New OWASP Dev Guide – Crypto chapter
- Blog: <http://off-the-wall-security.blogspot.com/>
- GitHub: <https://github.com/kwwall/>
- Email: <kevin.w.wall@gmail.com>
- Twitter: @KevinWWall
- CISSP, GIAC Web Application Defender (GWEB)





# What I will cover

- Dev good news / bad news
- Mistakes in using the following:
  - Pseudo random number generators
  - Secure hashes
  - Symmetric encryption
  - Asymmetric encryption
- Miscellaneous topics (time permitting)
  - Key management
  - Transparent DB Encryption
  - Java SSLSocket
  - JCE crypto providers



# Good News / Bad News

- Good news:

- Devs no longer designing their own crypto
- Devs rarely implementing standard algorithms

- Bad news:

- ✗ Dev “expertise” from copy-&-paste from Stack Overflow, etc., so still get things wrong.
  - ✗ Confidentiality vs. authenticity
  - ✗ Confusion of cipher modes, padding schemes
- ✗ Broken crypto for legacy applications
- ✗ Even experts still get things wrong (e.g., OpenSSL, GPG, etc.).

# Pseudo Random Number Generators (PRNG)

# PRNG Weaknesses

- Having a good source of (pseudo) randomness is essential to good cryptography.
  - Poor randomness → broken crypto
  - Cryptographers demand a “cryptographically secure” PRNG (CSRNG)
    - `java.util.Random` is *not* a CSRNG
    - `java.security.SecureRandom` is a CSRNG
  - CSRNG must have unpredictable seed
    - Seed entropy must equal (and should exceed) the internal state of the CSRNG



# PRNG Weaknesses: What to look for

- Using `java.util.Random` for *anything* related to crypto—this would include keys, IVs, nonces, etc.
- Seeding any CSRNG with insufficient entropy
  - If you initially require N-bits of randomness, then the entropy pool should have *at least* N-bits of randomness.
  - Generally not a problem with the default Oracle/Sun implementation of `SecureRandom` and `SHA1PRNG`.
    - Default `SecureRandom` CTOR uses `/dev/urandom` when available ***BUT*** may a problem if lots of randomness is required at boot time or if no `/dev/urandom` or `/dev/random`



# Example of correct use / seeding of SecureRandom

```
SecureRandom csrng =  
    SecureRandom.getInstance("SHA1PRNG", "BC");  
  
csrng.setSeed(  
    csrng.generateSeed( 160/8 )  
);    // Call before nextBytes(), etc.
```

For JDK 8 and later, consider using  
    SecureRandom.getInstanceStrong()  
instead of SecureRandom.getInstance().



# Secure Cryptographic Hashing

# Secure Hashing Weaknesses:

## What to look for (1/4)

- Use of trivially broken algorithms: MD2, MD4, MD5 or algorithms that are not true message digests such as CRCs.
- Use of less easily compromised broken algorithms: SHA1 (may be okay for legacy use for backward compatibility and some CSRNG cases).



# Secure Hashing Weaknesses:

## What to look for (2/4)

- If concerned about *local* attacks...
  - Time-dependent comparison of hashes
    - E.g., Bad: `String.equals()` or `Arrays.equals()`
    - `MessageDigest.isEqual()` is okay *after* JDK 1.6.0\_17
- Calling `MessageDigest.digest(byte[])` or `update(byte[])` methods on unbounded input under adversary's control. (DoS attack)



# Secure Hashing Weaknesses:

## What to look for (3/4)

- Misusing secure hash (MessageDigest) for message authentication codes (MAC):
  - MAC is a *keyed* hash, where the key is a secret key generally shared out-of-band.
  - Incorrect, naïve use:

$$\text{MAC}(\text{key}, \text{message}) := \text{H}(\text{key} || \text{message})$$

Where '||' is bitwise concatenation.

**Problem:** Susceptible to “length extension attacks”.

- Correct use: Use an HMAC (RFC 2104)...

```
Mac hmac = Mac.getInstance("HmacSHA256", "SunJCE");  
hmac.init(key);
```



# Secure Hashing Weaknesses:

## What to look for (4/4)

- Misusing a secure hash to mask data where enumeration of all or most of the input space is feasible.
  - E.g., Use SHA-256(SSN) to store as key in database or to track in log file.
  - Problem: If adversary can observe hashes, she can enumerate SHA-256 hashes of all possible SSNs and compare these to stored hashes.



# Is use of MD5 *ever* okay?

- Best collision attack against it is now about  $O(2^{24.1})$ , which takes at most 5 or 6 seconds on a modern desktop / laptop.
- But...okay in following cases:
  - Used as a PRNG when we only need something that is more or less unique and unpredictable; example IV generation used with CBC for symmetric ciphers.
  - Used as an HMAC construct as defined in RFC 2104
    - Bellare, Canetti & Krawczyk (1996): Proved HMAC security doesn't require that the underlying hash function be collision resistant, but only that it acts as a pseudo-random function.

# Symmetric Encryption



# Symmetric Encryption Weaknesses

- Inappropriate cipher algorithms
  - You aren't *still* using RC4, are you?
- Insufficient key size: < 128 bits
  - Java: DESede defaults to 2-key TDES (112-bit) unless the JCE Unlimited Strength Jurisdiction Policy files are installed.
- “ASCII” generated keys
- Inappropriate use of cipher modes
  - Related: IV abuses
- Assuming confidentiality implies data integrity.





# (Non-extended) ASCII Keys

- Keys generated from passwords or passphrases.

E.g.,

```
String key = "#s0meSeCR3tK3y!!"; // Or from property
```

```
SecretKeySpec skey =
```

```
    new SecretKeySpec( key.getBytes(), "AES");
```

```
Cipher cipher =
```

```
    Cipher.getInstance("AES/CBC/PKCS5Padding");
```

```
cipher.init(Cipher.ENCRYPT_MODE, skey);
```

```
...
```

ASCII: 96 printable bit codes vs 256 bits for random key



# Inappropriate use of cipher modes

**Question:** `Cipher.getInstance( "AES" ) ...`

What's the default cipher mode?

- Block modes and stream modes
  - Block modes: ECB and CBC
  - Stream modes: pretty much everything else
- All modes except for ECB require an IV.
- Streaming modes: Must **not** reuse the same key / IV pair... ***EVER!***
- Streaming modes do not require padding.



# Inappropriate use of cipher modes:

## ECB

- ECB is the raw application of the cipher algorithm.
- Reasons why it is the most commonly misused:
  - First (and sometimes only) example in textbooks
  - Simplest to implement (no need to bother with IVs)
- Weaknesses:
  - Same plaintext blocks always encrypt to same ciphertext
  - Block replay attacks are possible

# What's Wrong with ECB Mode?


**Original Tux  
image**



**Tux image  
encrypted with  
ECB mode**



**Tux image  
encrypted with  
any other cipher  
mode**



From:

via Wikimedia Commons; Larry Ewing, [lewing@isc.tamu.edu](mailto:lewing@isc.tamu.edu), and The GIMP.



# ECB: Block Replay Attack (1/6)

- Adversary can modify encrypted message without knowing the key or even encryption algorithm.
  - Can mangle message beyond recognition.
    - Remove, duplicate, and/or interchange blocks
  - Can usurp meaning of message if structure known. Consider the following scenario...



# ECB: Block Replay Attack (2/6)

[Example from Schneier, *Applied Cryptography*]\*

- Assume 8-byte encryption block size.
- Money transfer system to move money between banks
- Assume bank's standard message format is:

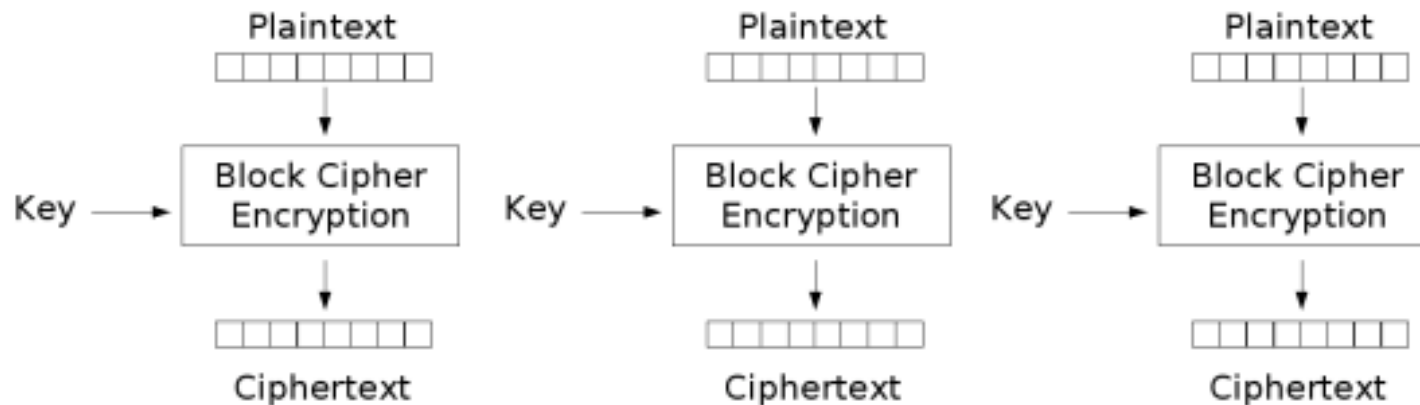
Bank 1: Sending**	1.5	blocks
Bank 2: Receiving**	1.5	blocks
Recipient's Name	6	blocks
Recipient's Acct#	2	blocks
Deposit Amount	1	block

-----

\* First discussed by C. Campbell, *IEEE Computer*, 1978

\*\* Sending / Receiving would be bank's "routing #s" today.

# ECB: Block Replay Attack (3/6)



Electronic Codebook (ECB) mode encryption

Each block is encrypted (and decrypted) independently

# ECB: Block Replay Attack (4/6)

- Mallory is MITM agent, listening to the comm channel between Bank of Alice and Bank of Bob.
- Mallory sets up accounts in both banks and deposits seed money in Bank of Alice.
- Mallory transfers some fixed amount of the seed money from Bank of Alice to her account at the Bank of Bob and records the transaction.
- Repeats later, and looks for identical blocks; eventually isolates her encrypted acct transfer information.





# ECB: Block Replay Attack (5/6)

- Mallory can now insert those message blocks into communication channel at will. Each time, that fixed amount will be deposited in Mallory's account at the Bank of Bob.
  - Mallory could also just inject her encrypted name and acct # for any transactions to the Bank of Bob, but why wait for such transactions?
- Two banks will notice by close of business when accts are reconciled. By that time, Mallory has already skipped town.



# ECB: Block Replay Attack (6/6)

- Can *not* be defeated by simply prepending date/time stamp to bank transfer authorization message. Mallory can replay individual blocks that lie on whole block boundaries (e.g., in this case the Recipient's Name and account #).
- *Can* be defeated by adding secure *keyed* hash to entire message (or using another cipher mode).



# ECB: What to look for

- No cipher mode specified at all. E.g.,  
`Cipher cipher = Cipher.getInstance("AES");`  
In Java, this is the same as:  
`Cipher cipher =  
 Cipher.getInstance("AES/ECB/PKCS5Padding");`
- No evidence that an IV is used
  - In Java, look for *absence* of both `IVParameterSpec` and `Cipher.getIV()`
  - Check lengths of resulting encryption
    - Generally IV is prepended to the raw ciphertext.  
(Exception might be where IV is fixed (bad) or determined algorithmically; discussed later.)



# ECB: Is it ever okay?

- Yes, when:
  - Encrypting plaintext with a less than 1 cipher block and ciphertext attacks not feasible:
    - Blowfish and DES (and hence DESede) block size: 64 bits
    - AES block size (and most other AES candidates): 128 bits
  - OR when encrypting *random* data
    - E.g., nonces, session IDs, *random* secret keys; maybe passwords if strong passwords enforced (LOL!).
- AND padding is used when appropriate (random data)
- AND block replay attacks are not an issue
- OR, using it for asymmetric encryption (only applicable mode!)




# If use of ECB *seems* okay...

- Make sure it is not used in a scenario where a block replay attack is possible.
- Ask yourself:
  - Are multiple blocks of ciphertext encrypted with ECB used?
  - Are these multiple ciphertext blocks exposed to an “adversary”?
  - Will block re-ordering ever fail to be detected in any cases? (I.e., are there cases where data integrity not always ensured?)
- If answer to these is “yes” for all questions, block replay is probably possible.

# Key / IV reuse in streaming mode (1/9)

- Stream ciphers and block ciphers operating in streaming modes create a cipher bit stream that is XOR'd with the plaintext stream.
- For a given key / IV pair, the same cipher bit stream is generated each time. Let's call this cipher bit stream,  $C(K, IV)$ .
- Let the encryption function for such a streaming mode be designated as  $E(K, IV, msg)$ .

Then,  $E(K, IV, msg) = msg \text{ XOR } C(K, IV)$



# Key / IV reuse in streaming mode (2/9)

- Let's see what happens if we encrypt 2 different plaintext messages, A and B, this way
- If an adversary intercepted both of these ciphertext results, they can compute the XOR of them, which is

$$-E(K, IV, A) = A \text{ XOR } C(K, IV)$$

$$-E(K, IV, B) = B \text{ XOR } C(K, IV)$$

$$E(K, IV, A) \text{ XOR } E(K, IV, B) =$$

$$A \text{ XOR } C(K, IV) \text{ XOR } B \text{ XOR } C(K, IV)$$

which, since XOR is commutative, is:

$$A \text{ XOR } B \text{ XOR } C(K, IV) \text{ XOR } C(K, IV) = A \text{ XOR } B$$

*That is, the XOR of the 2 plaintext messages, A and B.*



# Key / IV reuse in streaming mode (3/9)

- So what do we do with the XOR of 2 plaintext messages, A and B?
- If messages A and B are both written in some normal language (or character set, like ASCII), we can make that as a guess and use frequency distribution of some anticipated language (or format, such as C#s, etc.) and guess likely plaintext bits (characters). If the result resembles something intelligible (e.g., ASCII letter), guess was probably right.
- Modest computers can crack this in matter of few minutes for modest length messages.







# Key / IV reuse in streaming mode (4/9)

- The more ciphertexts created using the same key / IV pair and observed by an adversary, the better.
- Fixed message formats / structures (e.g., knowing you have all numeric fields such as SSN or credit card #) make it even more trivial.
- Eventually, both plaintexts (or shortest part if different lengths) get revealed.


# Key / IV reuse in streaming mode (5/9)

Next 4 slides from Dr. Rick Smith, Univ of St. Thomas, MN  
(License: Creative Commons Attribution ShareAlike 3.0 USA)




# Key / IV reuse in streaming mode (6/9)

- To recover the original message (image), we XOR the encrypted “Send Cash” image with the encryption key again:




# Key / IV reuse in streaming mode (7/9)



Note that we have the **same** encryption key XOR'ing both images.

# Key / IV reuse in streaming mode (8/9)


- Here's what happens when we XOR the 2 images that both used the same encryption key together:




"Send Cash"  
Encrypted



xor




Smiley Encrypted



yields



Both messages overlaid



# Key / IV reuse in streaming mode (9/9)

- *But wait!* It gets worse. If an application is doing this and an adversary can decrypt a message, they may be able to use a MITM attack to actually *alter* the ciphertext.
- Wikipedia example (Stream\_cipher\_attack):

$$\begin{aligned} (C(K) \text{ xor } "\$1000.00") \text{ xor } (" \$1000.00" \text{ xor } " \$9500.00") &= \\ C(K) \text{ xor } "\$1000.00" \text{ xor } "\$1000.00" \text{ xor } "\$9500.00" &= \\ C(K) \text{ xor } "\$9500.00" \end{aligned}$$



# Detour: Authenticated Encryption

- Encryption provides confidentiality, not integrity.  
(Integrity  $\neq$  authenticity)
- Approaches to authenticated encryption
  - **Encrypt-then-MAC (EtM)**: Encrypt, then apply MAC over IV+ciphertext and append the MAC.
  - **Encrypt-and-MAC (E&M)**: Encrypt the plaintext and append a MAC of the plaintext.
  - **MAC-then-Encrypt (MtE)**: Append a MAC of the plaintext and encrypt them both together.
- Decryption operation applied in reverse order.
- **EtM** built into some cipher modes such as CCM, GCM, EAX, etc.



# Horton Principle

- David Wagner and Bruce Schneier
- Relevant when considering what to data to include in a MAC
- Semantic authentication: “Authenticate what is meant, not what is said”
  - Avoid unauthenticated data: either don’t send / rely on it, or include it in the MAC
  - Relevant in message formats and protocols
- E.g., Alice sends:  
“metadata | | IV | | ciphertext | | MAC”





# Symmetric Encryption Weaknesses: CBC (when AE is not available)

- Overall, CBC probably most robust mode when used correctly.
- Use correctly means:
  - Random key and random IV with padding
  - HMAC over the IV+ciphertext applied as “encrypt-then-MAC” (EtM) approach.
- Common mistakes:
  - Fixed IV or predictable IV (e.g., counter, time, etc.)
  - Failure to MAC correctly (e.g., no MAC at all, encrypt-and-MAC, or MAC-then-encrypt)
  - Ignoring Horton principle (ESAPI 2.x!!!)

# Why is AE needed?

- When ciphertext's authenticity is in doubt, certain cryptographic attacks are possible that will either divulge the plaintext (or portions thereof) or possibly even reveal the secret key.
- Padding oracle attack, Serge Vaudenay, 2002
  - Originally discussed as deficiency in IPSec and SSL
  - Dismissed as being impractical until Rizzo and Duong research and POET software in 2010



# Symmetric Encryption Weaknesses:

Assuming confidentiality implies data integrity

- Only true if one is using an AE cipher mode such as CCM or GCM (the only 2 AE modes that are NIST approved and Java supported) or using a correctly implemented EtM approach.
- If confidentiality is not required, better (and faster) to just use an HMAC.
- Look for cases where plaintext is already known to attacker and encryption is used to prevent tampering.

# Asymmetric Cryptography: Encryption

# Asymmetric Ciphers and Chosen Plaintext Attacks (1/3)

- All asymmetric ciphers are prone to chosen plaintext attacks (CPA).
  - CPA is a cryptanalytic attack where an attacker can choose which plaintext to encrypt and then observe the resulting ciphertext.
  - CPA is always possible with asymmetric ciphers because we assume the algorithm details is known as well as the *public* key.

# Asymmetric Ciphers and Chosen Plaintext Attacks (2/3)

- Why might this be a problem?
  - Normally it's not because we usually are encrypting highly unpredictable plaintext that is too large to be enumerated.
    - E.g., symmetric session keys, cryptographic hash values
  - It becomes a problem when the plaintext is highly regular or short enough to enumerate all possible values



# Asymmetric Ciphers and Chosen Plaintext Attacks (3/3)

Real-life (bad) example:

- Application uses RSA algorithm to encrypt credit-card #s and store the resulting ciphertexts in application DB.
- Consider inside attacker with access to DB records (e.g., DBA, developer, tester) as well as the *public* key.
- Attacker encrypts all possible credit card #s with public key and saves mapping of plaintext / ciphertext pairs.
- Lookup into application DB records via CC# ciphertext allows discovery of credit card holder as well as revealing plaintext CC#.



# Miscellaneous Topics

- Key Management
- Database Encryption
- TLS/SSL issues in Java
- JCE Providers





# Key Management:

## Re-keying Frequency(1/2)

- PCI DSS 2.0 and later says that you *must* change symmetric crypto keys *at least* yearly? Is that enough?
- NIST SP 800-67, Revision 1 recommends:
  - For 3DES in CBC mode, re-key at least every  $2^{32} * 64$ -bits of plaintext
  - For AES in CBC mode, every  $2^{64} * 128$ -bits
  - General: every  $2^{N/2} * \text{cipher\_block\_size}$  bits, where N is key size in bits.
- New NIST 3DES recommendation (Rev 2): Lowered to every  $2^{20}$  blocks (i.e., every 8GB !!!)



# Key Management:

## Re-keying Frequency (2/2)

- “Sweet32”, a TLS attack on legacy 32-bit cipher suites is example:
  - <https://sweet32.info/>
  - [https://sweet32.info/SWEET32\\_CCS16.pdf](https://sweet32.info/SWEET32_CCS16.pdf)
- Matthew Green blog post provides more explanation:  
<https://blog.cryptographyengineering.com/2016/08/24/attack-of-week-64-bit-ciphers-in-tls/>



# Key Management: Secure Key Storage

*So where do you store your keys?*

- Ideally: an HSM or a TPM
- FAIL: If hard-coded in source code or put into properties file.
  - Both situations usually under version control!
- Ok: Config file, locked down & controlled by ops staff and unavailable to all others.
- Better: For .NET, DPAPI, WebLogic Encryption Services, Java Key Store
- NEVER put encryption key in *same* file with data that's being encrypted.



# Encrypting Data in a DB

*Three ways to encrypt data for a database:*

1. DB Engine itself does it via (mostly) Transparent Data Encryption (TDE)
  2. Done via a proxy; e.g., MIT's CryptDB
  3. Done via application code
- From application perspective, TDE approach is simplest.
  - Transparent to the application.
  - Available for Oracle and Microsoft SQL Server
  - Probably satisfies “letter of the law” for PCI DSS compliance (not verified).



# 30k' view of TDE

- Offers encryption at the column, table, and tablespace levels.
- Limited ciphersuite available; e.g., AES & 3DES
- Key management: usually 2 keys involved:
  - DB “master” key – a key encryption key, secured w/ password
  - Table / column / tablespace keys, encrypted by DB master key
- Usually CBC mode used, with usually with same IV for all encryptions
  - Same IV required for deterministic encryption so indexing works as expected
  - “Salt” allows non-deterministic encryption



# WIYTM? Why TDE fails

- If any application has that DB table / column open, then any other application with access to that table / column has access to encrypted data!
  - Not problem if data properly partitioned via “views”.
  - Backups, depending on how done, can be in plaintext!
- Usually the data we are encrypting in DB is:
  - Less than 20 bytes
  - Has particular format
  - Limited possible values

Result: Patterns may allow enumeration of values.



# SSLSocket & Server AuthN

- SSLSocket (or subclass) created by SSLSocketFactory does not do host name verification or cert pinning by default. Hence, MITM attacks are possible.
  - Must implement your own. 2 approaches:
    - Subclass SSLSocket; see <http://www.velocityreviews.com/forums/t958287-adding-hostname-verification-to-sslsocket.html>
    - Create an SSLContext that does host name verification; see <http://stackoverflow.com/questions/8545685/writing-a-ssl-checker-using-java>

# Specifying JCE Providers

- Java has a concept of security providers.
  - Statically added via:
    - JRE: `$JAVA_HOME/lib/security/java.security`
    - JDK: `$JAVA_HOME/jre/lib/security/java.security`
  - Dynamically added via:
    - `Security.addProvider(Provider provider)`
    - `Security.insertProviderAt(Provider provider, int pos)`
    - Various `getInstance()` methods take `Provider` as 2<sup>nd</sup> arg
- Determined by position; defaults to what is in `java.security`.
- This concept extends to crypto providers





# What could possibly go wrong?

```
import org.bouncycastle.jce.provider.*;
```

```
...
```

```
int pos = Security.addProvider(  
    new BouncyCastleProvider() );
```



# Static setting in java.security

- Default list of providers ordered by preference:  
security.provider.1=sun.security.provider.Sun  
security.provider.2=sun.security.rsa.SunRsaSign  
security.provider.3=sun.security.ec.SunEC  
...  
security.provider.9=sun.security.smartcardio.SunPCSC  
security.provider.10=sun.security.mscapi.SunMSCAPI  
security.provider.11=org.bouncycastle.jce.provider.BouncyCastlePr  
ovider



# How about this?

```
import org.bouncycastle.jce.provider.*;
```

```
...
```

```
Security.insertProviderAt(  
    new BouncyCastleProvider(), 1 );
```



# Equivalent static setting in java.security

- Equivalent as if we did this:

security.provider.1=org.bouncycastle.jce.provider.Bouncy  
CastleProvider

security.provider.2=sun.security.provider.Sun

security.provider.3=sun.security.rsa.SunRsaSign

security.provider.4=sun.security.ec.SunEC

...

security.provider.10=sun.security.smartcardio.SunPCSC

security.provider.11=sun.security.mscapi.SunMSCAPI



# What could possibly go wrong?

- Consider this in `Logger.getLogger()` method in *rogue* copy of `log4j.jar` someone downloaded:

...

```
Security.insertProviderAt(  
    new MyEvilProvider(), 1 );
```

...



# How do we address this?

- Specify the Provider instance as part of the `getInstance()` methods; e.g.,

```
Cipher.getInstance("AES/CBC/PKCS5Padding",  
                  new BouncyCastleProvider() );
```

OR

- Use a Java Security Manager and restrict what classes may call `Security.addProvider()` and `Security.insertProviderAt()`

# What to look for

- Calls to either  
    `Security.addProvider()`

OR

`Security.insertProviderAt()`

without the use of a Java Security Manager  
(JSM)

**Caveat:** Java Security Manager is rarely used and if it is used, usage of a properly restrictive security policy is hardly ever set. Also, if the jars are not signed and validated before use, using the JSM matters little.



# Additional References

- New OWASP Dev Guide, chapter 11 (Cryptography) [still a work in progress]
  - <https://github.com/OWASP/DevGuide/blob/master/03-Build/0x11-Cryptography.md>
  - And those references therein



# Questions?

Now, after the conference, or email me at  
[kevin.w.wall@gmail.com](mailto:kevin.w.wall@gmail.com)

Or DM me on Twitter: @KevinWWall  
(must follow first)